

# Unified Shared Memory

March 2020

# What is Unified Shared Memory

Unified Shared Memory (USM) is:

- A pointer-based approach for data management in DPC++
- Complements buffers, not a replacement


Why USM?

- Simplify porting existing C++ codes to DPC++
- Give desired level of control over data movement

Why is it called USM?

- OpenMP calls it that – don't reinvent the wheel


# USM in a Picture



# Allocation Types

Three types of allocations:

- Device
  - Accessible on device, not on host
  - “Give me a pointer to my GPU’s DRAM”
- Host
  - Accessible on host and device
  - Allocated in pinned Host memory, does not migrate to device DRAM
- Shared
  - Accessible on host and device
  - Can migrate between host and device DRAM



# Buffer Example


Declare C++ Arrays

```
auto A = (int *) malloc(N * sizeof(int));
auto B = (int *) malloc(N * sizeof(int));
auto C = (int *) malloc(N * sizeof(int));

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

{
    buffer<int, 1> Ab(A, range<1>{N});
    buffer<int, 1> Bb(B, range<1>{N});
    buffer<int, 1> Cb(C, range<1>{N});

    q.submit([&] (handler& h) {
        auto R = range<1>{N};
        auto aA = Ab.get_access<access::mode::read>(h);
        auto aB = Bb.get_access<access::mode::read>(h);
        auto aC = Cb.get_access<access::mode::write>(h);
        h.parallel_for(R, [=] (id<1> i) {
            aC[i] = aA[i] + aB[i];
        });
    });
    q.wait();
} // A,B,C updated
```



Declare C++ Arrays


Initialize C++ Arrays

```
auto A = (int *) malloc(N * sizeof(int));
auto B = (int *) malloc(N * sizeof(int));
auto C = (int *) malloc(N * sizeof(int));

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

{
    buffer<int, 1> Ab(A, range<1>{N});
    buffer<int, 1> Bb(B, range<1>{N});
    buffer<int, 1> Cb(C, range<1>{N});

    q.submit([&] (handler& h) {
        auto R = range<1>{N};
        auto aA = Ab.get_access<access::mode::read>(h);
        auto aB = Bb.get_access<access::mode::read>(h);
        auto aC = Cb.get_access<access::mode::write>(h);
        h.parallel_for(R, [=] (id<1> i) {
            aC[i] = aA[i] + aB[i];
        });
    });
    q.wait();
} // A,B,C updated
```



Declare C++ Arrays

Initialize C++ Arrays


Declare Buffers

```
auto A = (int *) malloc(N * sizeof(int));
auto B = (int *) malloc(N * sizeof(int));
auto C = (int *) malloc(N * sizeof(int));

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

{
    buffer<int, 1> Ab(A, range<1>{N});
    buffer<int, 1> Bb(B, range<1>{N});
    buffer<int, 1> Cb(C, range<1>{N});

    q.submit([&] (handler& h) {
        auto R = range<1>{N};
        auto aA = Ab.get_access<access::mode::read>(h);
        auto aB = Bb.get_access<access::mode::read>(h);
        auto aC = Cb.get_access<access::mode::write>(h);
        h.parallel_for(R, [=] (id<1> i) {
            aC[i] = aA[i] + aB[i];
        });
    });
    q.wait();
} // A,B,C updated
```



Declare C++ Arrays

Initialize C++ Arrays

Declare Buffers

Declare Accessors

```
auto A = (int *) malloc(N * sizeof(int));
auto B = (int *) malloc(N * sizeof(int));
auto C = (int *) malloc(N * sizeof(int));

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

{
    buffer<int, 1> Ab(A, range<1>{N});
    buffer<int, 1> Bb(B, range<1>{N});
    buffer<int, 1> Cb(C, range<1>{N});

    q.submit([&] (handler& h) {
        auto R = range<1>{N};
        auto aA = Ab.get_access<access::mode::read>(h);
        auto aB = Bb.get_access<access::mode::read>(h);
        auto aC = Cb.get_access<access::mode::write>(h);
        h.parallel_for(R, [=] (id<1> i) {
            aC[i] = aA[i] + aB[i];
        });
    });
    q.wait();
} // A,B,C updated
```



Declare C++ Arrays

Initialize C++ Arrays

Declare Buffers

Declare Accessors


Use Accessors in Kernel

```
auto A = (int *) malloc(N * sizeof(int));
auto B = (int *) malloc(N * sizeof(int));
auto C = (int *) malloc(N * sizeof(int));

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

{
    buffer<int, 1> Ab(A, range<1>{N});
    buffer<int, 1> Bb(B, range<1>{N});
    buffer<int, 1> Cb(C, range<1>{N});

    q.submit([&] (handler& h) {
        auto R = range<1>{N};
        auto aA = Ab.get_access<access::mode::read>(h);
        auto aB = Bb.get_access<access::mode::read>(h);
        auto aC = Cb.get_access<access::mode::write>(h);
        h.parallel_for(R, [=] (id<1> i) {
            aC[i] = aA[i] + aB[i];
        });
    });
    q.wait();
} // A,B,C updated
```



Declare C++ Arrays

Initialize C++ Arrays

Declare Buffers

Declare Accessors

Use Accessors in Kernel

C++ Arrays Updated

```
auto A = (int *) malloc(N * sizeof(int));
auto B = (int *) malloc(N * sizeof(int));
auto C = (int *) malloc(N * sizeof(int));

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

{
    buffer<int, 1> Ab(A, range<1>{N});
    buffer<int, 1> Bb(B, range<1>{N});
    buffer<int, 1> Cb(C, range<1>{N});

    q.submit([&] (handler& h) {
        auto R = range<1>{N};
        auto aA = Ab.get_access<access::mode::read>(h);
        auto aB = Bb.get_access<access::mode::read>(h);
        auto aC = Cb.get_access<access::mode::write>(h);
        h.parallel_for(R, [=] (id<1> i) {
            aC[i] = aA[i] + aB[i];
        });
    });
    q.wait();
} // A,B,C updated
```

# USM Example

Declare USM Arrays

```
int *A = malloc_shared<int>(N, q);
int *B = malloc_shared<int>(N, q);
int *C = malloc_shared<int>(N, q);

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

q.submit([&] (handler& h) {
    auto R = range<1>{N};
    h.parallel_for(R, [=] (id<1> ID) {
        auto i = ID[0];
        C[i] = A[i] + B[i];
    });
});
q.wait();
// A,B,C updated and ready to use
```

Declare USM Arrays

Initialize USM Arrays

```
int *A = malloc_shared<int>(N, q);
int *B = malloc_shared<int>(N, q);
int *C = malloc_shared<int>(N, q);

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

q.submit([&] (handler& h) {
    auto R = range<1>{N};
    h.parallel_for(R, [=] (id<1> ID) {
        auto i = ID[0];
        C[i] = A[i] + B[i];
    });
});
q.wait();
// A,B,C updated and ready to use
```

Declare USM Arrays

Initialize USM Arrays

Read/Write USM Arrays

```
int *A = malloc_shared<int>(N, q);
int *B = malloc_shared<int>(N, q);
int *C = malloc_shared<int>(N, q);

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

q.submit([&] (handler& h) {
    auto R = range<1>{N};
    h.parallel_for(R, [=] (id<1> ID) {
        auto i = ID[0];
        C[i] = A[i] + B[i];
    });
});
q.wait();
// A,B,C updated and ready to use
```

Declare USM Arrays

Initialize USM Arrays

Read/Write USM Arrays

USM Arrays Updated

```
int *A = malloc_shared<int>(N, q);
int *B = malloc_shared<int>(N, q);
int *C = malloc_shared<int>(N, q);

for (int i = 0; i < N; i++) {
    A[i] = i; B[i] = 2*i;
}

q.submit([&] (handler& h) {
    auto R = range<1>{N};
    h.parallel_for(R, [=] (id<1> ID) {
        auto i = ID[0];
        C[i] = A[i] + B[i];
    });
});
q.wait();
// A,B,C updated and ready to use
```

# Allocation Routines

- Many flavors defined:
  - `void*` returning `malloc_shared`, `malloc_device`, `malloc_host`
  - `T` returning `malloc_shared<T>`, `malloc_device<T>`, `malloc_host<T>`
  - `void*` returning `malloc(..., kind)`
  - `T` returning `malloc<T>(..., kind)`
  - C++ allocator (useful with containers)
- Aligned versions exist as well
- Allocations must take a context (and sometimes a device)
  - Can also just pass a queue and use its context/device
  - Device and shared allocations need to allocate against a specific device
  - Allocations are not guaranteed to be usable across different contexts
- Free takes a context or queue.

# Explicit Data Movement

Device allocations cannot be directly accessed on the host

- Programmers must manually copy data between host and device
- memcpy(...) on queue or handler classes
  - Currently async by default
- memset(...) currently, fill(...) in future
  - byte vs T

Trades off extra complexity for full control over data movement

Host allocations have no data movement.



# Implicit Data Movement

Shared allocations do not require programmers to manually transfer data between host and device.

- Data movement will be automatically handled by some combination of HW, drivers, and low-level runtimes.

Trades off control (and performance) for simplicity and productivity

*However:*

- Programmers can give automatic systems additional information to change their behavior
- Prefetch, Memadvise, etc.

# Shared Allocations: Restricted or not?

USM defines two capability levels for shared allocations:

- Restricted
- Concurrent

Restricted:

- Shared allocs are basically device allocs visible on host
- Limited to device memory, but still migrates between host and device
- Should not concurrently access allocations on host and device

Concurrent:

- Shared allocs are just allocs in the “shared” address space
- Migrates freely and different parts of the allocation can be concurrently accessed on host and device (Think page granularity)
- Not limited to device memory

# Task Scheduling

## Explicit Scheduling

- Submitting a kernel returns an Event
- Wait on Events to order tasks

```
auto E = q.submit([&] (handler& h) {  
    auto R = range<1>{N};  
    h.parallel_for(R, [=] (id<1> ID) {  
        auto i = ID[0];  
        C[i] = A[i] + B[i];  
    });  
});  
E.wait();
```

## DPC++ Graph Scheduling

- Build Task Graphs from Events

```
auto R = range<1>{N};  
  
auto E = q.submit([&] (handler& h) {  
    h.parallel_for(R, [=] (id<1> ID) {...});  
});  
  
q.submit([&] (handler& h) {  
    h.depends_on(E);  
    h.parallel_for(R, [=] (id<1> ID) {...});  
});
```

# Device Queries

- Are device allocations supported?
- Are host allocations supported?
- Are shared allocations supported?
- Are shared allocations restricted?
- Is the system allocator supported?

# Pointer Queries

Two flavors:

- Given a ptr, what type of USM allocation is it, if any?
- Given a ptr, what device was it allocated against, if it's a shared/device allocation?

# History: Why not OpenCL SVM?

OpenCL defined 3 flavors of Shared Virtual Memory (SVM):

- Coarse-grain buffer – Too hard to use
- Fine-grain buffer – Too hard to use
- System – Too hard to implement

We have also proposed USM for OpenCL

# Summary

USM is for pointer-based data management in DPC++

- Multiple types of allocations
- Explicit and Implicit data movement
- Event-based task scheduling

Enables other API simplifications to reduce DPC++ verbosity

- `q.parallel_for`, etc.