

oneAPI Data Parallel C++ Library

For the DPC++ Technical Advisory Board discussion

May 2020

Agenda

1. Top-level namespace & include directory for oneAPI (cont.)
2. oneDPL execution policies
3. Algorithms: Synchronous vs. Asynchronous
4. Algorithms: Range-based API
5. (if time remains) Extension API overview

1. Top-level namespace for oneAPI

- Following the recommendation from the previous TAB meeting, we plan to introduce a common top-level namespace for oneAPI libraries
- Two related questions we would like to hear your feedback on:
 - a) The preferred name for the namespace
 - b) The preferred choice of top-level include directories

Naming options

- `namespace oneapi { ... }`
 - Obvious correlation with oneAPI
 - Less chances of collision with other APIs
- `namespace one { ... }`
 - Short and easy to read
 - Together with next-level (library) namespace, better matches the library short names and APIs using those (e.g. version macros)

```
oneapi::std::for_each(oneapi::dpl::execution::gpu, oneapi::dpl::begin(buf), oneapi::dpl::end(buf),  
    [...](...){ /*code here*/ });
```

```
one::std::for_each(one::dpl::execution::gpu, one::dpl::begin(buf), one::dpl::end(buf),  
    [...](...){ /*code here*/ });
```

```
using namespace oneapi; // or one  
std::for_each(dpl::execution::gpu, dpl::begin(buf), dpl::end(buf),  
    [...](...){ /*code here*/ });
```

Top level include directory

- The include structure often mirrors the namespace structure
- **Do you see value in *identically named* top level include directories?**

```
#include "dpl/algorithm"  
#include "tbb/parallel_for.h"  
#include "mkl/blas/blas.hpp"
```

```
#include "onedpl/algorithm"  
#include "onetbb/parallel_for.h" // backward incompatible  
#include "onemkl/blas/blas.hpp"
```

```
#include "oneapi/dpl/algorithm"  
#include "oneapi/tbb/parallel_for.h" // backward incompat.  
#include "oneapi/mkl/blas/blas.hpp"  
// or  
#include "one/dpl/algorithm"  
#include "one/tbb/parallel_for.h" // backward incompat.  
#include "one/mkl/blas/blas.hpp"
```

```
onedpl/major.minor/  
  include/  
    oneapi/  
      dpl/  
        algorithm, ...  
  test/  
  ...  
onetbb/major.minor/  
  include/  
    oneapi/  
      tbb -> ../tbb  
  tbb/  
    parallel_for.h, ...  
src/  
...
```

2. oneDPL execution policies

```
template <typename KernelName = /*unspecified*/>
class device_policy //: public parallel_unsequenced_policy
{
public:
    device_policy();
    explicit device_policy( sycl::queue queue );
    explicit device_policy( sycl::device device );
    template <typename OtherName>
    device_policy( const device_policy<OtherName>& policy );

    sycl::queue queue() const; // also considering implicit conversion
};
```

Examples:

```
using namespace one::dpl::execution;
device_policy<class my_kernel_name> policy{sycl::gpu_selector{}};
auto pol = make_device_policy(sycl::queue{});
```

Predefined execution policies

- Idea: create predefined policy objects for commonly used devices
 - similar to `std::execution::{par, seq, unseq, par_unseq}`
- **Which way of naming do you like more?**

Verbose (4 variations)

```
namespace one::dpl::execution {  
    inline device_policy<> default_policy; //1  
    ... cpu_policy_v {sycl::cpu_selector{}}; //2  
    ... gpu_pol      {sycl::gpu_selector{}}; //3  
    ... policy_gpu    {sycl::gpu_selector{}}; //4  
}  
...  
std::reduce(default_policy, ...); //1  
std::for_each(cpu_policy_v, ...); //2  
std::sort(dpl::execution::gpu_pol, ...); //3
```

Concise (similar to the standard)

```
namespace one::dpl::execution {  
    inline device_policy<> deflt; // maybe dev?  
    ... cpu {sycl::cpu_selector{}};  
    ... gpu {sycl::gpu_selector{}};  
}  
...  
std::reduce(deflt, ...);  
std::for_each(cpu, ...);  
std::sort(dpl::execution::gpu, ...);
```

3. Synchronous vs. asynchronous

- Now oneDPL algorithms with DPC++ execution policies vary in behavior
- Synchronous: the function waits until execution completes on the device
 - Standard-compliant; in some cases, may transfer the data back to the host
 - Used for function that return a value: `reduce`, `find`, ...
- Implicitly Asynchronous: the function submits a kernel and returns
 - Non-standard; requires explicit data transfer or waiting on the queue
 - More efficient: allows SYCL to build kernel graphs/pipelines
 - Used for functions with no return value: `for_each`, `transform`, `sort`, ...
- Long-term plans are to add explicitly asynchronous APIs
 - Thrust went through this: v1.9.4 made a breaking change for all “std-like” calls to `block`, and introduced asynchronous algorithms

Options for oneDPL specification v1.0

- a) Describe the current behavior: specify some algorithms as blocking and some as non-blocking
- b) Allow either blocking or non-blocking implementation, except where the semantics requires a certain behavior
- c) Require all algorithms to block, compliant with C++ standard
 - Makes the initial oneDPL release not compliant OR losing performance
- d) Same as c), plus define asynchronous APIs
 - Risks making it wrong due to lack of implementation/prototype expertise

Which option does seem right to you?

4. Range-based API for algorithms

- C++20 adds Ranges into the C++ standard library
 - Very powerful and expressive functional API
 - But does not yet support execution policies
- We work on adding range support for oneDPL algorithms
 - Only for a subset of the standard algorithms and views
 - Not fully standard-compliant (not based on concepts, no projections, ...)

Ranges: programmability and kernel fusion

A pipeline of 3 kernels:

```
std::reverse(pol, begin(data), end(data));  
std::transform(pol, begin(data), end(data), begin(result), lambda1);  
auto res = std::find_if(pol, begin(result), end(result), pred);
```

With fancy iterators (1 kernel):

```
auto res = std::find_if(pol,  
    make_transform_iterator(make_reverse_iterator(end(data)), lambda1),  
    make_transform_iterator(make_reverse_iterator(begin(data)), lambda1),  
    pred);
```

With ranges (1 kernel):

```
auto res = one::dpl::find_if(pol,  
    views::all(data) | views::reverse() | views::transform(lambda1), pred);
```

Ranges: planned scope (23 algorithms)

<algorithm>

for_each	
find	copy
find_if_not	transform
find_if	sort
search	stable_sort
search_n	partial_sort
min_element	partial_sort_copy
max_element	is_sorted_until
minmax_element	is_sorted

<numeric>

reduce
transform_reduce
exclusive_scan
transform_exclusive_scan
inclusive_scan
transform_inclusive_scan

The set of range views is being defined

Range-based API for algorithms (summary)

- C++20 adds Ranges into the C++ standard library
 - Very powerful and expressive functional API
 - But not yet for algorithms with execution policies
- We work on adding range support for oneDPL algorithms
 - Only for a subset of the standard algorithms and views
 - Not fully standard-compliant (not based on concepts, no projections, ...)
 - **Are any important algorithms missed in the minimal subset?**
- For the oneDPL specification, the options are:
 - a) Add these APIs to the oneDPL specification v1.0, or
 - b) Leave it to a later version (i.e. implement as extensions in Intel's oneDPL)
 - **Which option does seem right to you?**

5. oneDPL extension APIs (as of now)

reduce_by_segment	Partial reductions on a sequence of values, by segments of equal keys
inclusive_scan_by_segment	Partial prefix scans on a sequence of values, by segments of equal keys
exclusive_scan_by_segment	
binary_search	Binary search variations for multiple values in the same sequence ("vectorized" search)
lower_bound	
upper_bound	
identity, maximum, minimum	functional utility classes
zip_iterator	Iterates over multiple sequences simultaneously
transform_iterator	Applies a function to each element in a sequence
permutation_iterator	Iterates over a sequence of values in the order set by a sequence of indices
counting_iterator	"Iterates" over a virtual sequence of numbers (holds a counter)
begin, end	Functions to pass a SYCL buffer to parallel algorithms (return an object holding a position in the buffer)

The principles of adding new extension APIs

- Functionality: the API is required to support a desired use case
- Complexity: desired API semantics does not allow a “trivial” mapping to existing APIs
- Performance: trivial mapping to existing APIs is not performant, and optimizations make it non-trivial
- Convenience: meaningful and commonly used API name (even with a simple mapping)
- Consistency: significant semantical similarity with APIs selected by other criteria
- Explicit demand with a reasonable justification