

# oneMKL Technical Advisory Board

Session 2

June 3, 2020

# Agenda

- Welcoming remarks – 5 minutes
- Overview of oneMKL programming model – Maria Kraynyuk (25 minutes)
- Walk-thru of the oneMKL Specification – Spencer Patty (20 minutes)
- Wrap-up and next steps – 5 minutes

# oneMKL TAB Members

- Mike Dewar, Numerical Algorithms Group (NAG)
- Mark Hoemmen, Stellar Science
- Nevin Liber, Argonne National Laboratory (ANL)
- Piotr Luszczek, Innovative Computing Laboratory (ICL) at University of Tennessee, Knoxville (UTK)
- Pat Quillen, MathWorks
- Nichols Romero, ANL
- Harry Waugh, University of Bristol

# oneAPI Math Kernel Library Programming Model

# Example: Matrix Multiplication (1 of 3)

## Buffer API

```
int64_t n = 32;

sycl::device dev(sycl::{host,cpu,gpu}_selector());
sycl::queue Q(dev);

double *A, *B, *C;
A = (double *)malloc(n*n*sizeof(double));
...

sycl::buffer<double, 1> A_buf{A, range<1>(n * n)},
                        B_buf{B, range<1>(n * n)},
                        C_buf{C, range<1>(n * n)};

onemkl::blas::gemm(Q,
                   onemkl::transpose::N, onemkl::transpose::N,
                   n, n, n, 1.0,
                   A_buf, n,
                   B_buf, n, 0.0,
                   C_buf, n);
```

device setup

prepare  
matrices

$C = A * B$

## C API

```
int64_t n = 32;

double *A, *B, *C;
A = (double *)malloc(n*n*sizeof(double));
...

cblas_dgemm(CblasRowMajor,
            CblasNoTrans, CblasNoTrans,
            n, n, n, 1.0,
            A, n,
            B, n, 0.0,
            C, n);
```

# Example: Matrix Multiplication (2 of 3)

## Buffer API

```
using onemkl::blas::gemm;
int64_t n = 32;
sycl::device dev(sycl::{host,cpu,gpu}_selector());
sycl::queue Q(dev);

double *A = ..., *B = ..., *C = ...;

sycl::buffer<double, 1> A_buf{A, range<1>(n * n)},
                        B_buf{B, range<1>(n * n)},
                        C_buf{C, range<1>(n * n)};

gemm(Q, onemkl::transpose::N,
      onemkl::transpose::N,
      n, n, n, 1.0, A_buf, n, B_buf, n,
      0.0, C_buf, n);
```

device setup

prepare  
matrices

$C = A * B$

## USM API

```
using onemkl::blas::gemm;
int64_t n = 32;
sycl::device dev(sycl::{host,cpu,gpu}_selector());
sycl::queue Q(dev);

size_t bytes = n * n * sizeof(double);
double *A = sycl::malloc_shared(bytes, dev,
                                Q.get_context());
double *B = sycl::malloc_shared(...);
double *C = sycl::malloc_shared(...);

gemm(Q, onemkl::transpose::N,
      onemkl::transpose::N,
      n, n, n, 1.0, A, n, B, n,
      0.0, C, n);

Q.wait_and_throw();
```

# Example: Matrix Multiplication (3 of 3)

## Buffer API

```
using onemkl::blas::gemm;
int64_t n = 32;
sycl::device dev(sycl::{host,cpu,gpu}_selector());
sycl::queue Q(dev);
```

```
double *A = ..., *B = ..., *C = ...;
```

```
sycl::buffer<double, 1> A_buf{A, range<1>(n * n)},
                        B_buf{B, range<1>(n * n)},
                        C_buf{C, range<1>(n * n)},
                        D_buf{D, range<1>(n * n)};
```

```
gemm(Q, onemkl::transpose::N,
      onemkl::transpose::N,
      n, n, n, 1.0, A_buf, n, B_buf, n,
      0.0, C_buf, n);
```

```
gemm(Q, onemkl::transpose::N,
      onemkl::transpose::N,
      n, n, n, 1.0, C_buf, n, A_buf, n,
      0.0, D_buf, n);
```

device setup

prepare  
matrices

$C = A * B$

$D = C * A$

## USM API

```
using onemkl::blas::gemm;
int64_t n = 32;
sycl::device dev(sycl::{host,cpu,gpu}_selector());
sycl::queue Q(dev);
```

```
size_t bytes = n * n * sizeof(double);
double *A = sycl::malloc_shared(bytes, dev,
                                Q.get_context());

double *B = sycl::malloc_shared(...);
double *C = sycl::malloc_shared(...);
double *D = sycl::malloc_shared(...);
```

```
sycl::event e =
gemm(Q, onemkl::transpose::N,
      onemkl::transpose::N,
      n, n, n, 1.0, A, n, B, n,
      0.0, C, n);
```

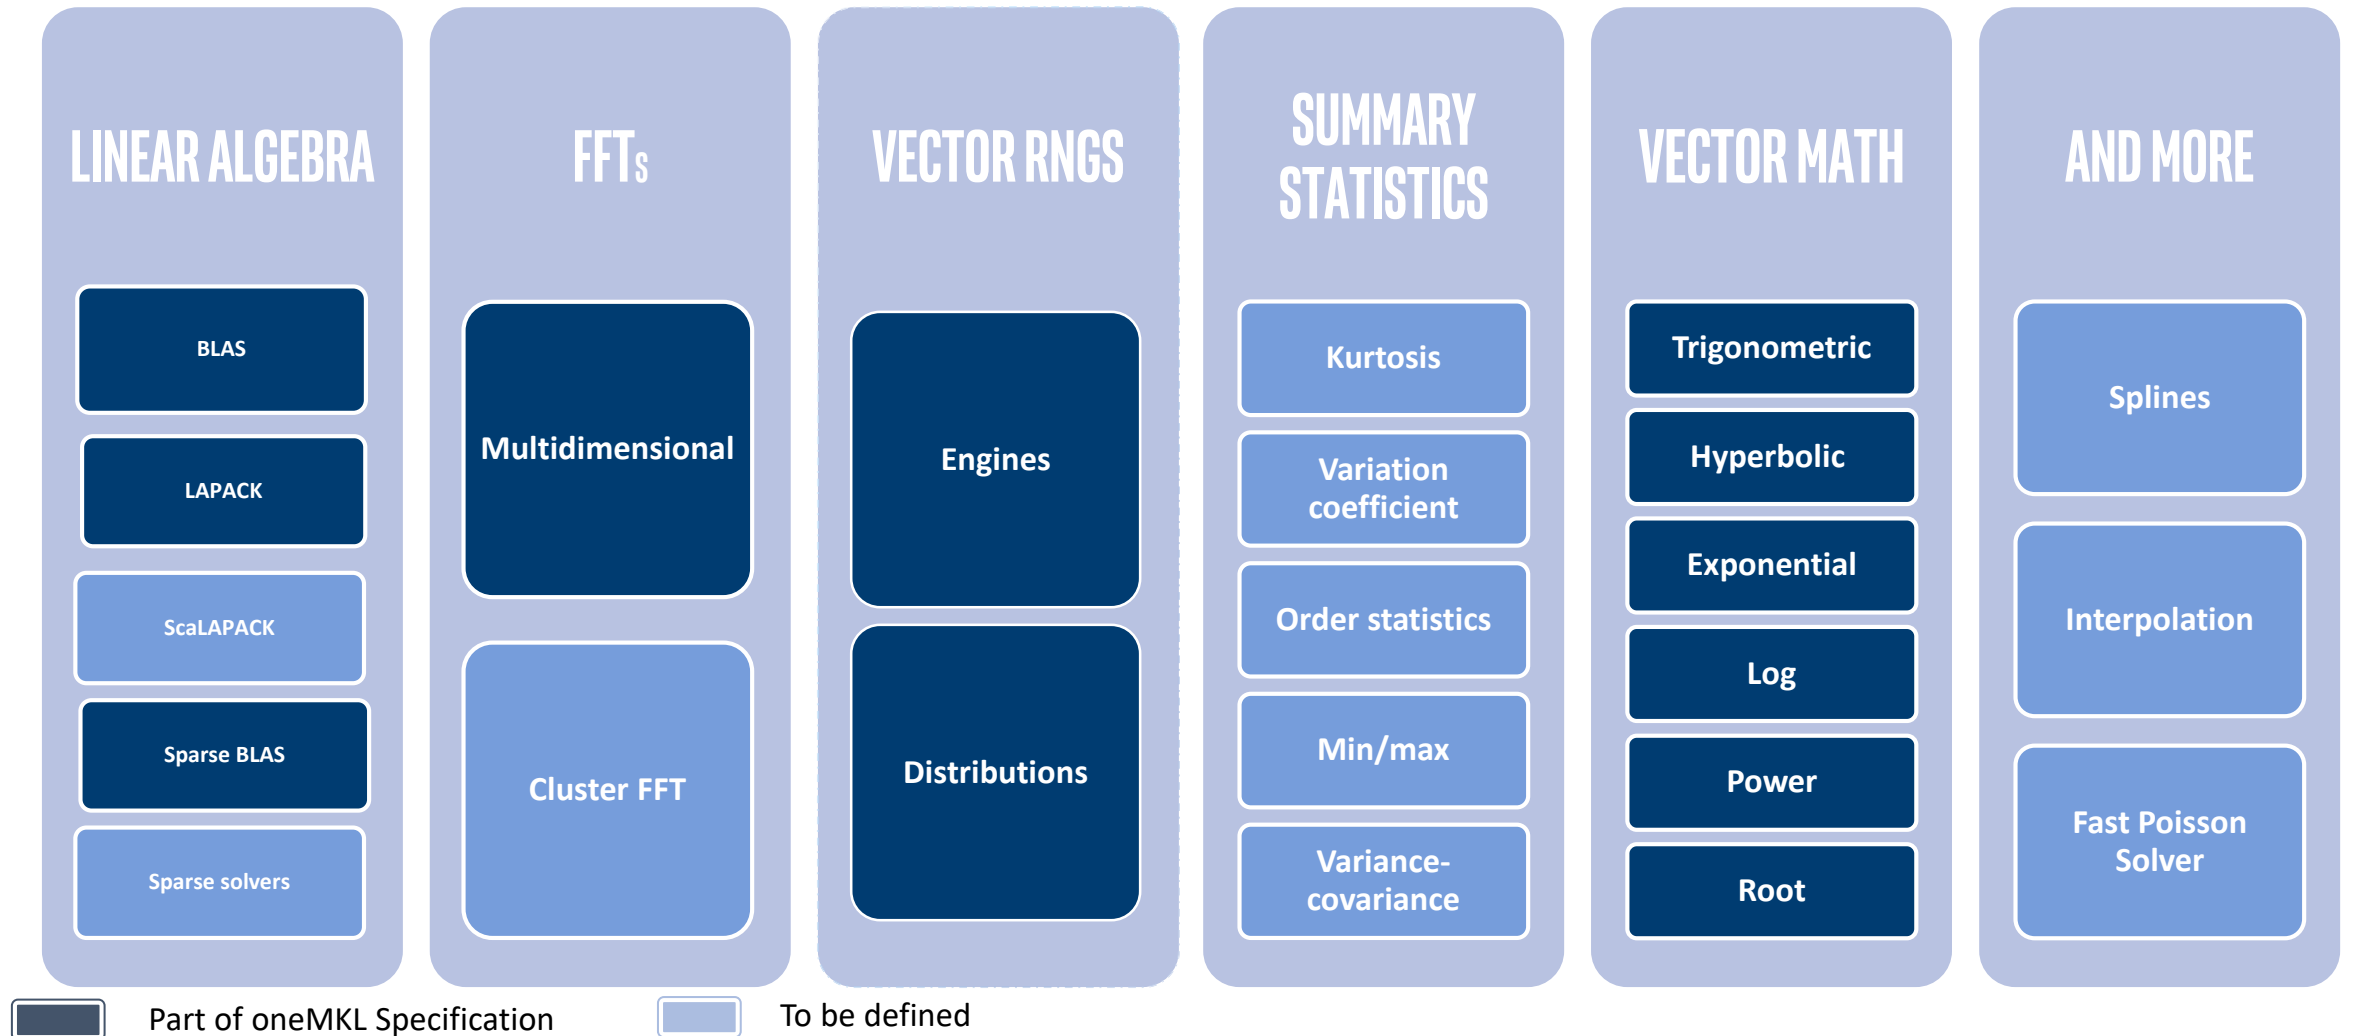
```
gemm(Q, onemkl::transpose::N,
      onemkl::transpose::N,
      n, n, n, 1.0, C, n, A, n,
      0.0, D, n, {e});
```

# Walk-thru of the oneMKL Specification

<https://spec.oneapi.com/versions/latest/elements/oneMKL/source/index.html>



# oneAPI Math Kernel Library Specification



# Next Steps

- Look over current oneMKL Spec v. 0.8
  - Focus most on oneMKL Architecture (section 1) and BLAS/LAPACK APIs
- Focuses for next meeting(s):
  - Overview of the open source oneMKL interfaces GitHub project
  - Particular feedback requests on oneMKL Spec:
    - Asynchronous execution
    - Multi GPU execution
    - Exceptions/error codes

# Resources

- oneAPI Main Page: <https://www.oneapi.com/>
- Latest release of oneMKL Spec (currently v. 0.8):  
<https://spec.oneapi.com/versions/latest/elements/oneMKL/source/index.html>
- GitHub for oneAPI Spec: <https://github.com/oneapi-src/oneAPI-spec>
- GitHub for oneAPI TAB: <https://github.com/oneapi-src/oneAPI-tab>
- Latest build of oneAPI Spec: <http://staging.spec.oneapi.com.s3-website-us-west-2.amazonaws.com/exclude/ci/branches/refs/heads/master/versions/latest/index.html>
- GitHub for open source oneMKL interfaces (currently BLAS domain):  
<https://github.com/oneapi-src/oneMKL>