

oneMKL Technical Advisory Board

Session 6

August 12, 2020

Agenda

- Welcoming remarks – 5 minutes
- Updates from last meeting – 10 minutes
- oneMKL Random Number Generators pass downs and open questions - Pavel Dyakov and Alina Elizarova (15 minutes)
- Overview of oneMKL Summary Statistics domain – Pavel Dyakov and Alina Elizarova (15 minutes)
- Wrap-up and next steps – 5 minutes

oneMKL Specification v. 0.9

- Released July 30!
- Modifications to oneMKL Architecture, BLAS and LAPACK domains
- Significant refactoring and updating of Sparse BLAS, VM, RNG, and DFT domains API descriptions and structure
- Add Summary Statistics domain
- Add future considerations and acknowledgment to appendices
- Change top-level namespace to `oneapi::mkl`

oneMKL Specification v. 1.0 and later

- Targeting v. 1.0 (release date August 30):
 - Exceptions
 - Minor RNG changes
 - Formatting, typo fixes
- Considered for future versions:
 - Encapsulation of matrix and vector objects
 - More human-readable names
 - Broader support for row major layout
 - Alternative handling of computational failures

oneMKL TAB Meeting Frequency

- Will present overviews of more domains:
 - Sparse linear algebra
 - Discrete Fourier transforms
 - Vector math
- Any requests for other topics?
- Change frequency to every 4 weeks?

oneMKL Random Number Generators Pass Downs and Open Questions

Actions Taken from Previous Meeting for oneMKL RNG

- All trivial constructors for distributions are removed from specification (0.9)
- Keep engines with copy constructor (deep-copy) and add move constructor (will be updated in 1.0)

oneMKL RNG Engines

- `oneapi::mkl::rng::mrg32k3a`
 - `oneapi::mkl::rng::philox4x32x10`
 - `oneapi::mkl::rng::mt19937`
 - `oneapi::mkl::rng::mt2203`
 - `oneapi::mkl::rng::mcg31m1`
 - `oneapi::mkl::rng::mcg59`
 - `oneapi::mkl::rng::ars5`
 - `oneapi::mkl::rng::nondeterministic`
 - `oneapi::mkl::rng::sfmt19937`
 - `oneapi::mkl::rng::whichmann_hill`
 - `oneapi::mkl::rng::r250`
 - `oneapi::mkl::rng::sobol`
 - `oneapi::mkl::rng::niederreiter`
- Reliable, performable, well-known and widely used
- Small period, but quite fast and usable
- Performance and/or implementation is HW dependent
- Popular before, but outdated
- Quasi-random

Open questions:

- Is "default_engine" needed?

Proposal: Adding "default_engine" to oneMKL spec as "implementation defined" (to support users who don't care about underlying engines)

- HW dependent engines: should we have such engines in oneMKL spec?

Proposal: Keeping HW dependent engines in oneMKL spec (to provide performable engines for the particular devices)

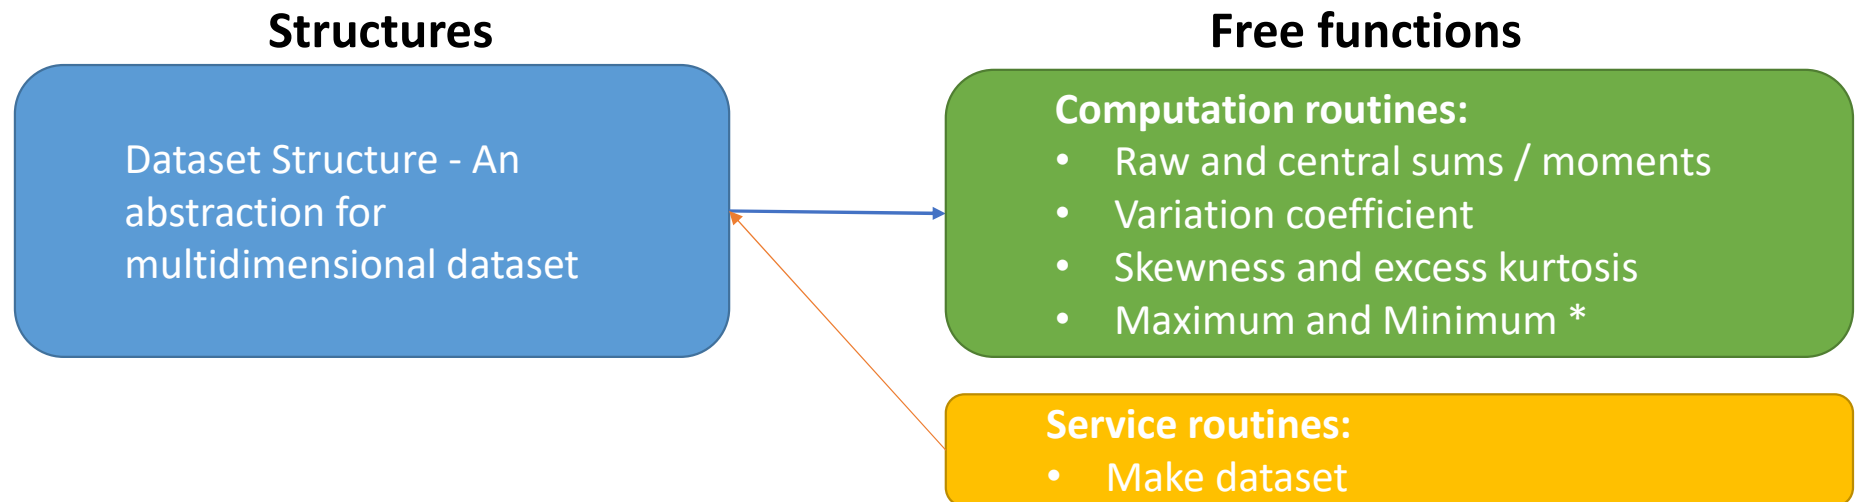
- Outdated engines: do we want to keep these engines in oneMKL spec?

Proposal: Keeping outdated engines in oneMKL spec (to provide wide set of engines)

Overview of oneMKL Summary Statistics Domain

oneMKL Summary Statistics Overview and Structure

- oneMKL Summary Statistics provide a set of routines that compute basic statistical estimates for single and double precision datasets.



* Routines set is targeted to be extended after 1.0 specification version

oneMKL Summary Statistics Usage Model

A typical algorithm for summary statistics is as follows:

- Create and initialize an object for dataset.
- Call the summary statistics routine to calculate the appropriate estimate.

* currently 1D buffers and USM pointers are used to represent multi-dimensional dataset

```
#include "CL/sycl.hpp"
#include "mkl_stats_sycl.hpp"

int main() {
    sycl::queue queue;

    const size_t n_observations = 1000;
    const size_t n_dims = 3;
    std::vector<float> x(n_observations * n_dims);
    // [fill x storage]

    //create buffer for dataset
    sycl::buffer<float, 1> x_buf(x.data(), x.size());
    // create buffer for mean values
    sycl::buffer<float, 1> mean_buf(n_dims);
    // create oneapi::mkl::stats::dataset
    auto dataset = oneapi::mkl::stats::make_dataset<oneapi::mkl::
stats::layout::row_major>(n_dims, n_observations, x_buf);

    oneapi::mkl::stats::mean(queue, dataset, mean_buf);

    //...
    return 0;
}
```

oneMKL Summary Statistics Dataset Structure

- Supported layout:
 - `oneapi::mkl::stats::row_major`
 - `oneapi::mkl::stats::col_major`

Parameter	Description
<code>n_dims</code>	the number of dimensions
<code>n_observations</code>	the number of observations
<code>observations</code>	the matrix of observations
<code>weights</code>	an optional parameter, represents array of weights for observations (of size <code>n_observations</code>). If the parameter is not specified, each observation is assigned a weight equal 1.
<code>indices</code>	an optional parameter, represents array of dimensions that are processed (of size <code>n_dims</code>). If the parameter is not specified, all dimensions are processed.

```
// Buffer-based
template<typename Type, layout ObservationsLayout>
struct dataset<sycl::buffer<Type, 1>, ObservationsLayout> {
    explicit dataset(std::int64_t n_dims_,
                    std::int64_t n_observations_,
                    sycl::buffer<Type, 1> observations_,
                    sycl::buffer<Type, 1> weights_ = {0},
                    sycl::buffer<std::int64_t, 1> indices_ = {0});

    std::int64_t n_dims;
    std::int64_t n_observations;
    layout = ObservationsLayout;
    sycl::buffer<Type, 1> observations;
    sycl::buffer<Type, 1> weights;
    sycl::buffer<std::int64_t, 1> indices;
};

// USM based
template<typename Type, layout ObservationsLayout>
struct dataset<Type*, ObservationsLayout> {
    explicit dataset(std::int64_t n_dims_,
                    std::int64_t n_observations_,
                    Type* observations_, Type* weights_ = nullptr,
                    std::int64_t* indices_ = nullptr);

    std::int64_t n_dims;
    std::int64_t n_observations;
    layout = ObservationsLayout;
    Type* observations;
    Type* weights;
    std::int64_t* indices;
};
```

oneMKL Summary Statistics Service Routines

Make Dataset

Service function to create a dataset object (for C++ version < 17 to avoid explicit USM/buffer template parameter instantiation)

Supported types:

- float
- double

// Buffer-based API. Make dataset

```
template<layout ObservationsLayout = row_major, typename Type>
dataset<sycl::buffer<Type, 1>, ObservationsLayout>
make_dataset(std::int64_t n_dims, std::int64_t n_observations,
             sycl::buffer<Type, 1> observations,
             sycl::buffer<Type, 1> weights = {0},
             sycl::buffer<std::int64_t, 1> indices = {0});
```

// USM-based API. Make dataset

```
template<layout ObservationsLayout = layout::row_major, typename Type>
dataset<Type*, ObservationsLayout>
make_dataset(std::int64_t n_dims, std::int64_t n_observations,
             Type* observations,
             Type* weights = nullptr,
             std::int64_t* indices = nullptr);
```

oneMKL Summary Statistics Routines

Mean / Skewness / Variation / Kurtosis / Min / Max

Compute mean / skewness* /
variation* / kurtosis* / min / max
values

Supported types:

- float
- double

Supported methods:

- oneapi::mkl::stats::method::fast
- oneapi::mkl::stats::method::one_pass**

* Support user provided mean case

** Not supported for min / max and skewness / variation /
kurtosis/ with user provided mean

// Buffer-based API. Mean

```
template<method Method = method::fast, typename Type,  
layout ObservationsLayout>  
void mean(sycl::queue& queue,  
          const dataset<sycl::buffer<Type, 1>, ObservationsLayout>&  
          data,  
          sycl::buffer<Type, 1> mean);
```

// USM-based API. Mean

```
template<method Method = method::fast, typename Type,  
layout ObservationsLayout>  
sycl::event mean(sycl::queue& queue,  
                 const dataset<Type*, ObservationsLayout>& data,  
                 Type* mean,  
                 const sycl::vector_class<sycl::event> &dependencies = {});
```

oneMKL Summary Statistics Routines

Central Sums and Moments

Compute central sums (moments)
up to the 4th order

Supported types:

- float
- double

Supported methods:

- oneapi::mkl::stats::method::fast
- oneapi::mkl::stats::method::one_pass

// Buffer-based API. Central sums

```
template<method Method = method::fast, typename Type,  
layout ObservationsLayout>  
void raw_sum(sycl::queue& queue,  
             const dataset<sycl::buffer<Type, 1>, ObservationsLayout>&  
             data,  
             sycl::buffer<Type, 1> central_sum_2,  
             sycl::buffer<Type, 1> central_sum_3 = {0},  
             sycl::buffer<Type, 1> central_sum_4 = {0});
```

// USM-based API. Central sums

```
template<method Method = method::fast, typename Type,  
layout ObservationsLayout>  
sycl::event raw_sum(sycl::queue& queue,  
                   const dataset<Type*, ObservationsLayout>& data,  
                   Type* central_sum_2,  
                   Type* central_sum_3 = nullptr,  
                   Type* central_sum_4 = nullptr,  
                   const sycl::vector_class<sycl::event> &dependencies = {});
```

Next Steps

- Look over current oneMKL Spec v. 0.9
- Focuses for next meeting(s):
 - Sparse linear algebra
 - Discrete Fourier transforms
 - Vector math

Version	Date	oneAPI Notes	oneMKL Notes
0.8.5	26 June 2020	80% content	Finalize BLAS and LAPACK domains
0.9.0	30 Jul 2020	~100% content	Finalize FFT, sparse BLAS, RNG, and VM domains
1.0.0	30 Aug 2020	Gold Release	Minor cleanup

Resources

- oneAPI Main Page: <https://www.oneapi.com/>
- Latest release of oneMKL Spec (currently v. 0.9):
<https://spec.oneapi.com/versions/latest/elements/oneMKL/source/index.html>
- GitHub for oneAPI Spec: <https://github.com/oneapi-src/oneAPI-spec>
- GitHub for oneAPI TAB: <https://github.com/oneapi-src/oneAPI-tab>
- Latest build of oneAPI Spec: <http://staging.spec.oneapi.com.s3-website-us-west-2.amazonaws.com/exclude/ci/branches/refs/heads/master/versions/latest/index.html>
- GitHub for open source oneMKL interfaces (currently BLAS domain):
<https://github.com/oneapi-src/oneMKL>